

Name: Feiyu Zhang

Netid: Feiyuz2

Team name: 0 errors 0 warnings

School affiliation: ZJUI

Milestone 3

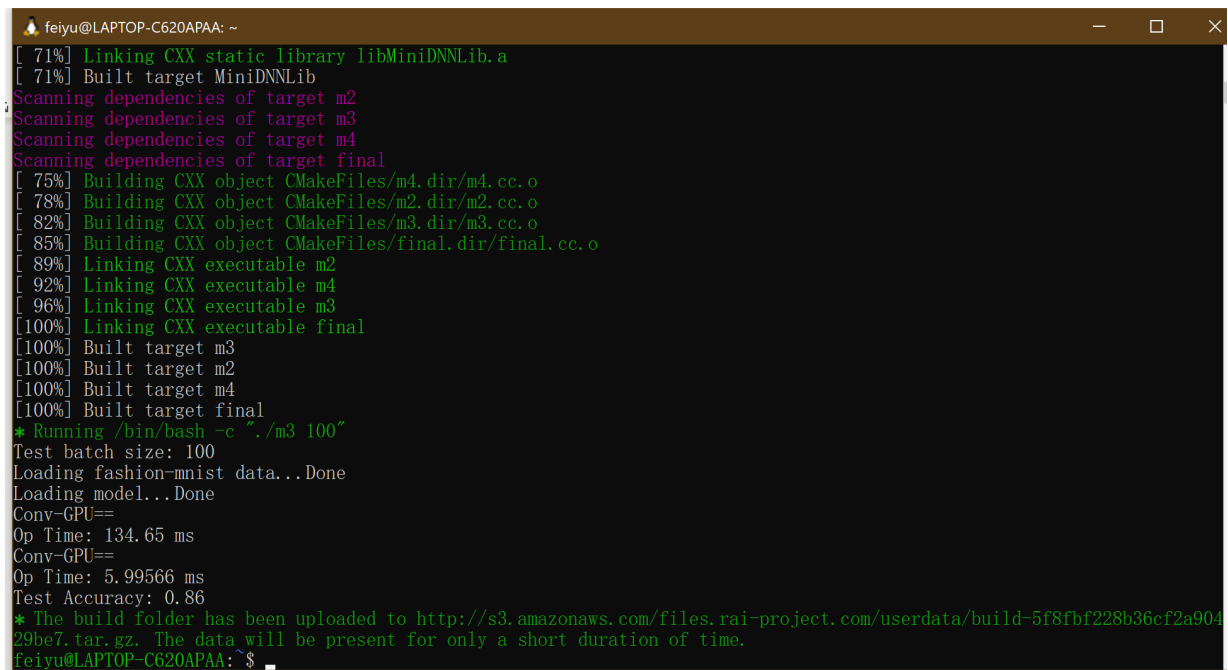
Deliverables

Everything from Milestone 2

Implement a GPU Convolution

Correctness and timing with 3 different dataset sizes

Report: Show output of rai running your GPU implementation of convolution



```
feiyu@LAPTOP-C620APAA: ~  
[ 71%] Linking CXX static library libMiniDNNLib.a  
[ 71%] Built target MiniDNNLib  
Scanning dependencies of target m2  
Scanning dependencies of target m3  
Scanning dependencies of target m4  
Scanning dependencies of target final  
[ 75%] Building CXX object CMakeFiles/m4.dir/m4.cc.o  
[ 78%] Building CXX object CMakeFiles/m2.dir/m2.cc.o  
[ 82%] Building CXX object CMakeFiles/m3.dir/m3.cc.o  
[ 85%] Building CXX object CMakeFiles/final.dir/final.cc.o  
[ 89%] Linking CXX executable m2  
[ 92%] Linking CXX executable m4  
[ 96%] Linking CXX executable m3  
[100%] Linking CXX executable final  
[100%] Built target m3  
[100%] Built target m2  
[100%] Built target m4  
[100%] Built target final  
* Running /bin/bash -c "./m3 100"  
Test batch size: 100  
Loading fashion-mnist data...Done  
Loading model...Done  
Conv-GPU==  
Op Time: 134.65 ms  
Conv-GPU==  
Op Time: 5.99566 ms  
Test Accuracy: 0.86  
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-5f8fbf228b36cf2a90429be7.tar.gz. The data will be present for only a short duration of time.  
feiyu@LAPTOP-C620APAA:~$
```

Deliverables

```
feiyu@LAPTOP-C620APAA: ~  
[ 71%] Linking CXX static library libMiniDNNLib.a  
[ 71%] Built target MiniDNNLib  
Scanning dependencies of target final  
Scanning dependencies of target m2  
Scanning dependencies of target m4  
Scanning dependencies of target m3  
[ 75%] Building CXX object CMakeFiles/final.dir/final.cc.o  
[ 78%] Building CXX object CMakeFiles/m2.dir/m2.cc.o  
[ 82%] Building CXX object CMakeFiles/m4.dir/m4.cc.o  
[ 85%] Building CXX object CMakeFiles/m3.dir/m3.cc.o  
[ 89%] Linking CXX executable m4  
[ 92%] Linking CXX executable m2  
[100%] Linking CXX executable final  
[100%] Linking CXX executable m3  
[100%] Built target m4  
[100%] Built target m2  
[100%] Built target m3  
[100%] Built target final  
* Running /bin/bash -c './m3 1000'  
Test batch size: 1000  
Loading fashion-mnist data...Done  
Loading model...Done  
Conv-GPU==  
Op Time: 281.088 ms  
Conv-GPU==  
Op Time: 54.9976 ms  
Test Accuracy: 0.886  
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-5f8fbf8d8b36cf2acb0f0d2b.tar.gz. The data will be present for only a short duration of time.  
feiyu@LAPTOP-C620APAA:~$  
  
feiyu@LAPTOP-C620APAA: ~  
Scanning dependencies of target m3  
Scanning dependencies of target m4  
Scanning dependencies of target m2  
Scanning dependencies of target final  
[ 75%] Building CXX object CMakeFiles/m4.dir/m4.cc.o  
[ 78%] Building CXX object CMakeFiles/m3.dir/m3.cc.o  
[ 82%] Building CXX object CMakeFiles/m2.dir/m2.cc.o  
[ 85%] Building CXX object CMakeFiles/final.dir/final.cc.o  
[ 89%] Linking CXX executable m2  
[ 92%] Linking CXX executable m3  
[ 96%] Linking CXX executable m4  
[100%] Linking CXX executable final  
[100%] Built target m2  
[100%] Built target final  
[100%] Built target m3  
[100%] Built target m4  
* Running /bin/bash -c './m3'  
Test batch size: 10000  
Loading fashion-mnist data...Done  
Loading model...Done  
Conv-GPU==  
Op Time: 803.28 ms  
Conv-GPU==  
Op Time: 522.156 ms  
Test Accuracy: 0.8714  
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-5f8f078c8b36cf04e2567297.tar.gz. The data will be present for only a short duration of time.  
feiyu@LAPTOP-C620APAA:~$
```

Report: Demonstrate nsys profiling the GPU execution

Deliverables

```
feiyu@LAPTOP-C620APAA: ~
CUDA Memory Operation Statistics (KiB)
-----
Total      Operations      Average      Minimum      Maximum      Name
-----
1722500.0      2      861250.0      722500.000      1000000.0      [CUDA memcpy DtoH]
538919.0      4      134729.0      0.766      288906.0      [CUDA memcpy HtoD]
Generating Operating System Runtime API Statistics...
Operating System Runtime API Statistics (nanoseconds)
-----
Time(%)      Total Time      Calls      Average      Minimum      Maximum      Name
-----
33.3      92147664263      935      98553651.6      51654      100194563      sem_timedwait
33.3      92085005495      934      98592083.0      60422      100281024      poll
21.7      60112526276      2      30056263138.0      22174371762      37938154514      pthread_cond_wait
11.6      32008966241      64      500140097.5      500089999      500179661      pthread_cond_timedwait
0.0      74480523      764      97487.6      1058      16736276      ioctl
0.0      19383466      9071      2136.9      1139      8245      read
0.0      3262991      97      33639.1      1073      1441787      mmap
0.0      616810      97      6358.9      1574      32041      open64
0.0      533336      1      533336.0      533336      533336      pthread_mutex_lock
0.0      224218      5      44843.6      32346      53785      pthread_create
0.0      75112      15      5007.5      1479      25732      munmap
0.0      71689      3      23896.3      10998      48587      fgetc
0.0      67180      15      4478.7      2189      9756      write
0.0      58888      3      19629.3      3426      36231      fopen64
0.0      58259      19      3066.3      1282      9835      fopen
0.0      41636      7      5948.0      2852      8356      fflush
0.0      27125      5      5425.0      2003      7764      open
0.0      16623      3      5541.0      4847      6868      pipe2
0.0      15208      8      1901.0      1000      5339      fclose
0.0      13807      2      6903.5      3746      10061      pthread_cond_signal
0.0      9843      2      4921.5      3976      5867      socket
0.0      5518      1      5518.0      5518      5518      fwrite
0.0      5317      1      5317.0      5317      5317      connect
0.0      1668      1      1668.0      1668      1668      bind
0.0      1397      1      1397.0      1397      1397      fcntl
Generating NVTX Push-Pop Range Statistics...
NVTX Push-Pop Range Statistics (nanoseconds)
* The build folder has been uploaded to http://s3.amazonaws.com/files.raai-project.com/userdata/build-5f8f08b18b36cf05432ca113.tar.gz. The d
t for only a short duration of time.
feiyu@LAPTOP-C620APAA: ~
Exported successfully to
/build/report1.sqlite
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)
-----
Time(%)      Total Time      Calls      Average      Minimum      Maximum      Name
-----
87.7      1109518392      6      184919732.0      82062      554069265      cudaMemcpy
12.1      152727642      6      25454607.0      61788      149885419      cudaMalloc
0.2      2224340      6      370723.3      56845      896359      cudaFree
0.0      253616      2      126808.0      32395      221221      cudaLaunchKernel
Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)
-----
Time(%)      Total Time      Instances      Average      Minimum      Maximum      Name
-----
100.0      116856830      2      58428415.0      23509740      93347090      conv_forward_kernel
CUDA Memory Operation Statistics (nanoseconds)
-----
Time(%)      Total Time      Operations      Average      Minimum      Maximum      Name
-----
92.7      915329284      2      457664642.0      385550060      529779224      [CUDA memcpy DtoH]
7.3      72165686      4      18041421.5      1568      38583992      [CUDA memcpy HtoD]
CUDA Memory Operation Statistics (KiB)
-----
Total      Operations      Average      Minimum      Maximum      Name
-----
1722500.0      2      861250.0      722500.000      1000000.0      [CUDA memcpy DtoH]
538919.0      4      134729.0      0.766      288906.0      [CUDA memcpy HtoD]
Generating Operating System Runtime API Statistics...
Operating System Runtime API Statistics (nanoseconds)
```

Deliverables

Report: Include a list of all kernels that collectively consume more than 90% of the program time.

```
Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)
```

Time (%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	116856830	2	58428415.0	23509740	93347090	conv_forward_kernel

Report: Include a list of all CUDA API calls that collectively consume more than 90% of the program time.

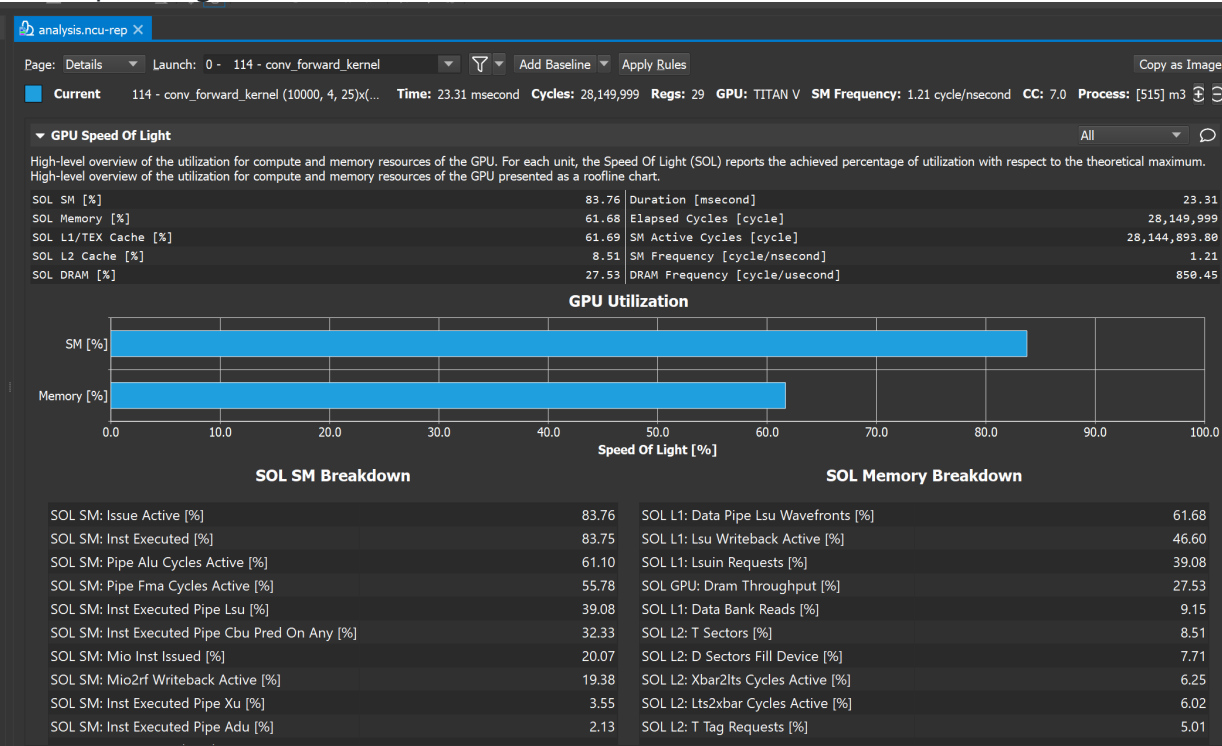
```
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)
```

Time (%)	Total Time	Calls	Average	Minimum	Maximum	Name
87.7	1109518392	6	184919732.0	82062	554069265	cudaMemcpy
12.1	152727642	6	25454607.0	61788	149885419	cudaMalloc
0.2	2224340	6	370723.3	56845	896359	cudaFree
0.0	253616	2	126808.0	32395	221221	cudaLaunchKernel

```
Generating CUDA Kernel Statistics...
Generating CUDA Memory Operation Statistics...
```

Report: Include an explanation of the difference between kernels and API calls
API calls are provided and kernels are written by us.
API calls are executed sequentially and kernels are executed parallelly.

Report: Screenshot of the GPU SOL utilization in Nsight-Compute GUI for your kernel profiling data



Deliverables

Use `rai -p <project folder> --queue rai_amd64_ece408 --submit=m3` to mark your job for grading

Milestone 2

Report: Show output of rai running Mini-DNN on the CPU (CPU convolution implemented) for batch size of 10k images

* Running `/bin/bash -c "time ./m2"`

Test batch size: 10000

Loading fashion-mnist data...Done

Loading model...Done

Conv-CPU==

Op Time: 84314.4 ms

Conv-CPU==

Op Time: 244010 ms

Test Accuracy: 0.8714

real 7m3.503s

user 7m2.510s

sys 0m0.992s

```
* Running /bin/bash -c "time ./m2"
Test batch size: 10000
Loading fashion-mnist data...Done
Loading model...Done
Conv-CPU==
Op Time: 84314.4 ms
Conv-CPU==
Op Time: 244010 ms
Test Accuracy: 0.8714
real    7m3.503s
user    7m2.510s
sys     0m0.992s
```

Report: List Op Times (CPU convolution implemented) for batch size of 10k images

Op Times:

Conv-CPU==

Op Time: 84314.4 ms

Conv-CPU==

Op Time: 244010 ms

Report: List whole program execution time (CPU convolution implemented) for batch size of 10k images

Whole execution time: user+sys=7m2.510s+0m0.992s=7m3.502s